

Analyzing the Semantic Modeling Capabilities of Google Sets

AZFAR KHANDOKER
Purdue University
akhandok@purdue.edu

RYAN MILLER
Purdue University
millerrv@purdue.edu

Abstract

In this paper we abstract data retrieved from Google Sets and transform the semantic concepts into a graph representation that we can further analyze. We use this representation to determine the potential of using Google Sets as a network building tool. We compare the generated graph with the graph representations of other online services, including Amazon co-purchasing and Wikipedia entries. We look for small world characteristics in all three networks. Finally, we show the similarities between these semantic networks using cosine similarity, and investigate the potential of using Google Sets as a product recommendation engine.

I. INTRODUCTION

There are many semantic networks publicly available. Amazon's co-purchasing dataset, which shows the products that are purchased together on amazon.com, and Wikipedia are a couple of the more interesting ones we examine in this paper. Google Sets, a public service provided by Google, associates a word or a list of words with a set of other words or lists of words, thus providing semantically similar results for a given query. All three of these services and their interactions therein can be abstracted into a graph representation, where semantic words or phrases make up the nodes in each network. In this paper, we describe how we transformed data retrieved from Google Sets into a graph that represents it as a semantic network. We also rigorously compare this network against the Amazon co-purchasing data set and the Wikipedia data set. We explore the possible implications of the similarities and the differences between these networks.

I.1 Semantic Networks

Popular in the fields of linguistics, psychology, and artificial intelligence, semantic networks are graphs used to illustrate the interconnected-

ness behind knowledge [1]. Although semantic networks can be structured to provide relationships (for example, hierarchy and inheritance), they can also model related concepts by connecting similar words or concepts with an edge. Our project explores characteristics behind semantic networks, such as the small world phenomenon and cosine similarity. In addition to established semantic networks, such as the Wikipedia online encyclopedia, we have implemented a strategy to build networks using data provided from Google Sets.

f_x | Indiana

	A	B
1	Indiana	
2		
3		
4		
5		

Figure 1

Google Sets data populates cells with semantically-related words after providing the root word “Indiana.”

f_x | Indiana

	A	B
1	Indiana	
2	indiana	
3	illinois	
4	michigan	
5	iowa	
6	florida	
7	ohio	
8	california	
9	minnesota	
10	georgia	
11	kansas	
12	wisconsin	
13	missouri	
14	kentucky	
15		

Figure 1
Google Sets data populates cells with semantically-related words after providing the root word "Indiana."

Table 1: Comparison semantic networks

	Wikipedia Snapshot	Amazon Co-purchasing
Number of nodes	5.7 million	542,000
Number of edges	130 million	1.2 million
Average degree	45.531	4.538
Average path length	Unavailable	2.842
Global Clustering Coefficient	0.014	0.215

I.2 Google Sets and Comparison Networks

Originally designed to add semantic language capability to search queries, Google Sets is a former Google Labs project released to the public in 2002¹. Despite being discontinued in 2011, data from Google Sets is still available through Google's cloud-based spreadsheet application. After entering a word into a spreadsheet cell, the user can hold the 'CTRL' or 'ALT' key and drag the bottom-right corner to cause semantically-related words to appear in the cells below the source word (Figure 1).

Google Sets provides a wealth of data derived from real-world searches that we aimed to utilize in our project, but we required a baseline to make sense of the semantic networks generated by Google Sets. Specifically, we compared networks created from Google Sets to other semantic networks, including the links among Wikipedia articles and product co-purchasing data from Amazon shopping (Table 1)^{2,3}. In the Wikipedia network, directed edges indicate that an article (encyclopedia entry) links directly to another article. Wikipedia editors can insert links arbitrarily into articles, leading to a well-connected network in which a Wikipedia article, on average, links to 45 other articles. In the Amazon co-purchasing network, nodes are represented by products (books, films, and others) available for sale on the Amazon website. When an Amazon

user purchases multiple items simultaneously, edges are added to the network to reflect "co-purchasing" conditions. Thus, according to our Amazon data set, an Amazon user, on average, purchases around 4 items together.

II. PROBLEM DEFINITION

First, we investigate whether the Google Sets data exhibits a small world phenomenon, especially in comparison to Wikipedia and Amazon co-purchasing networks. We then highlight the overlap between data from Google Sets and Amazon co-purchasing and determine whether the data can thus be used as a refinement tool for product recommendation engines. Google Sets, stemming from users' search queries, may therefore offer an ideal secondary source for recommendations. Finally, we address the issue of selecting the best candidate words from the network to recommend using the concept of cosine similarity, expanding on a study looking at semantic similarity of words and their synonyms appearing in the Roget's Thesaurus [2].

II.1 Related Work

Milgram was one of the first researchers to explore small world characteristics of a graph, noting the surprisingly short path lengths required to reach two non-neighbor individuals in a societal network [3]. He found a median of

¹<http://google.about.com/od/blogs/ss/Google-Labs-Dropouts-And-Failures.htm>

²Haselgrove, H. "Using the Wikipedia page-to-page link database." Retrieved from <http://haselgrove.id.au/wikipedia.htm>.

³J. Leskovec, L. Adamic and B. Adamic. (2007). "The Dynamics of Viral Marketing." *ACM Transactions on the Web (ACMTWEB)*. 1(1). Retrieved from <http://snap.stanford.edu/data/amazon-meta.html>

five hops required for two individuals to communicate (corresponding to an average path length of five, a bound we later observe in Google Sets data). Semantic networks in particular have also been shown to exhibit small world characteristics, such as English words and their synonyms, studied by [2].

Networks have also been used to generate recommendations, such as the use of membership statistics in social networks to predict movie recommendations [4]. Semantic networks can also aid in producing recommendations, such as a system that recommends music based on the semantic similarities of genres placed in a network [5].

II.2 Models, Measures and Algorithms

We use the metrics of average path length and degree distribution as described by Barabási and Albert [6] and clustering coefficient as described by Watts and Strogatz [7] to show whether a network is a small world network or not. As an important analysis tool used to further distinguish small world networks, we look for a power law degree distribution (Equation 1), rather than an exponential degree distribution, which does not indicate the same types of hubs present in small world networks.

$$p(x) \propto x^{-\alpha} \quad (1)$$

In addition to small world characteristic, we explore overlap of the networks with each other by comparing the results given by a traversal of each graph and counting the identical nodes encountered [8]. Finally, we use the notion of cosine similarity to show the similarity of the networks on a deeper semantic level compared with examining average path length and degree distribution [9]. Under the cosine similarity model (Equation 2), rows in the network's adjacency matrix for two nodes are represented as vectors. When two nodes share many of the same neighbors, their cosine similarity score

will be higher, indicating higher similarity between these two terms in semantic networks. As a product recommendation engine, these similar semantic terms will make ideal candidates for recommendation.

$$\sigma_{ij} = \cos \theta = \frac{n_{ij}}{\sqrt{d_i d_j}} \quad (2a)$$

$$n_{ij} = \sum_k A_{ik} A_{kj} \quad (2b)$$

III. IMPLEMENTATION

Our implementation consists of a Java application that we used to generate a graph representation of Google Sets. We also used various Python scripts to generate graphs of the Wikipedia dataset as well as the Amazon co-purchasing dataset. We also used Python scripts to run various analyses on the graphs.⁴

III.1 Google Sets Network Builder

Google Sets was only accessible via the Google Spreadsheet ("the spreadsheet") that is available as a part of Google Drive at the time of our studies. As such, a user would have to manually invoke the service by following the procedure we mentioned earlier. We needed to automate this process to allow for continued data retrieval even in the absence of human interaction. We created a robot to perform the user's actions using the Java libraries. We invoked the robot to perform these and other actions when we needed actions to be performed by the user. We also interfaced with the spreadsheet API to retrieve the results from Google Sets as well as fill in the spreadsheet with new data to be used with Google Sets. Finally, we used the *JGraphT*⁵ library to generate our graph.

⁴Our code repository can be found here: <http://code.google.com/p/590networks/>

⁵JGraphT website: <http://jgrapht.org/>




Figure 2
Subset of the star graph generated for the set of words produced in Figure 1, with "Indiana" serving as the center.

Every time we used Google Sets through the spreadsheet, we started with a seed word, s . For a seed word, we received a result set from Google Sets, $R = \{r_1, r_2, \dots, r_n\}$, of n result words where r_i denotes the i th result word. For every result word, we first checked if the result word matched a predefined regular expression. We did this to filter out any invalid words that we did not want to consider (e.g. words not in the English language or words containing numbers or symbols). The initial results for a seed word generate a star graph with the seed word as the central node (Figure 2).

III.2 Not Just a Collection of Stars

Next, we checked to see if r_i was the same as s for any i . If so, we ignored this result because we did not allow for self-loops in the graph. Finally, if the i th result was not already present in the graph, we created a node to represent it. If the result is already present in the graph, we did not create a new node for it, instead creating an edge between the node that represented r_i and the node that represented s and added it to the graph. Anytime we added a result word to the graph, we also added it to a queue of words that would give us the next s . We would repeat this procedure for all results in R .

One can think of Google Sets returning a star graph for each s . s would be the central node with each word in the returned R being the neighbors of it. However, since we only created a node for a word in R if and only if it was not already present in the graph, the structure of a star graph quickly dissipates since we would create edges between pre-existing nodes in the graph. For example, the result set may be $R_{\text{Indiana}} = \{\text{"Michigan"}, \text{"Iowa"}, \text{"Illinois"}\}$

for $s = \text{"Indiana"}$. This is a star graph. However, on the next iteration when we use $s = \text{"Michigan"}$, we may get a result set $R_{\text{Michigan}} = \{\text{"Illinois"}, \text{"Minnesota"}, \text{"Indiana"}\}$. We see that our graph is no longer a set of star graphs, since we have a triangle consisting of the nodes $\{\text{"Indiana"}, \text{"Michigan"}, \text{"Illinois"}\}$.

III.3 Analysis Tools

To analyze the graphs of Google Sets, Wikipedia and Amazon co-purchasing, we used the *igraph* library along with our own implementations of some analysis metrics. We relied on the standard *igraph* implementations of `degree_distribution()` and `path_length_hist()` when looking for small world characteristics. For the Wikipedia data set, at over 6 gigabytes in GML format, *igraph* was unable to return an average path length. Instead, we looked at various subgraphs to analyze this network property. When generating subgraphs at depth k from a source node, we rely on the `neighborhood()` procedure in *igraph*, which provides a useful way to examine subgraphs in large networks.

The *igraph* tool has no built-in procedure for cosine similarity, so we implemented this analysis technique in Python based on the formula provided in class (Equation 2). We also relied on standard graphing tools, such as Python's *matplotlib* and *Microsoft Excel* to analyze properties of the degree distribution curves and power law regression values.

IV. STUDIES

We conducted various studies on the graphs to see how the data contained in them may be similar and semantically related. We studied the small world characteristics of each of the networks and we also analyzed their semantic similarities.

IV.1 Small World Characteristics

To show the presence of small world phenomena, we analyzed the graphs using various met-

Table 2: Shortest paths from "wine" to "france"

Network	Path Length	Path
Google Sets	4	"wine" → "champagne" → "bordeaux" → "france"
Wiki	1	"Wine" → "France"
Amazon	4	"Wine" → "Italy (Culinar- naria)" → "Culinaria: The United States: A Culinary Discovery (Culinaria)" → "Culi- naria France (Culinaria Series)"

Table 3: Small world characteristics in comparison networks

	Google Sets	Wikipedia	Amazon
# nodes	2,871	544	1,373
# edges	8,962	25,403	3,432
Avg. degree	6.243	93.393	5.998
Avg. path length	4.329 (4.554)	2.344 (1.828)	4.287 (4.640)
Clustering coeff.	0.255 (0.002)	0.727 (0.172)	0.343 (0.003)

rics including average path length, clustering coefficient and degree distribution. When we analyzed the graphs, we ran a simple experiment on all three graphs to see how they each performed using the same input. We ran our algorithm with an initial seed of $s = \text{"wine"}$. We ran our algorithm until we either reached a depth of $k = 4$ or we found that one of the words in a returned R was the word "france". We summarize these findings in Table 2. This generated a subgraph for each of the three graphs. We used these subgraphs in our analyses. We summarize these subgraphs in Table 3. For the average path length and the clustering coefficient, we have also included some other data. The data in the parenthesis are the corresponding values for an Erdos-Renyi random graph [10] ("the random graph") with the same number of nodes and edges as the graph with which it is being compared (i.e. for the Google Sets subgraph, we created an Erdos-Renyi random graph with 2,871 nodes and 8,962 edges).

For average path length, only the Wikipedia data set demonstrated a higher average path

length than that of the random graph; the other two graphs show a smaller average path length. This uniqueness of Wikipedia among the other subgraphs is also apparent for the clustering coefficient. The clustering coefficients of Google Sets and Amazon are 128% and 114% higher than the random graph that coincides with them, respectively. However, for the Wikipedia data set, the clustering coefficient is merely 4% higher than that of the random graph for Wikipedia. Using these values, we can see that both Google Sets and Amazon demonstrate small world phenomenon characteristics but Wikipedia does not. Small average path length and large clustering coefficient (compared to random graphs of equal density) are necessary, but not sufficient conditions for a graph to be considered small world. We examined another characteristic, degree distribution, of the three graphs to further emphasize the presence, or absence, of small world phenomena in these graphs.

Figures 3, 4 and 5 show the degree distributions of the three subgraphs of Google Sets,

Wikipedia and Amazon, respectively. We have also included the power trendline equation for each of the graphs. We see that the equations for the Google Sets and Amazon subgraph demonstrate a power law degree distribution. For Equation 1 to be valid for a power law degree distribution, we expect an exponent value $2 < \alpha < 3$ [11]. Contrary to this, the Wikipedia data set does not have an α value within this range. We also see that the R^2 value for both Google Sets and Amazon are much higher than that for Wikipedia, indicating a strong power function curve fit. Such results are evident from examining the plots visually, as well. On a log-log scale, a power law degree distribution should resemble a line, which is shown by both Google Sets and Amazon. Wikipedia does not show the same behavior on a log-log scale plot. Thus it follows that a power law trend line was not a good fit for the degree distribution in the Wikipedia data set. We therefore conclude that the Wikipedia data set does not contain a power law degree distribution contrary to Google Sets and Amazon.




Figure 3: Google Sets - $p(x) = 0.174x^{-1.293}$, $R^2 = 0.724$




Figure 4: Wikipedia - $p(x) = 0.018x^{-0.341}$, $R^2 = 0.197$




Figure 5: Amazon - $p(x) = 0.878x^{-3.172}$, $R^2 = 0.964$

IV.2 Comparison of Semantic Networks

Comparing degree distribution and average path length give a global view of the entire network, but we also explore the actual words present in the network to see how the networks overlap, i.e. contain the same set of words. When we compare the nodes present in the Google Sets network for a seed word of "Oreo", we find high overlap with nodes present in the Amazon network, while this overlap is absent in comparison with the Wikipedia network (Tables 4 and 5).

Such a high overlap between Google Sets and Amazon data sets may imply a potential

Table 4: Semantic networks generated with a seed word of "Oreo"

	Google Sets	Wikipedia	Amazon
# nodes	10,229	16,478	1,119
# edges	33,081	1,252,227	1,118

Table 5: Node matching among the semantic networks

	Unique for Google Sets	Not in Google Sets	Shared	Percentage Shared
Google Sets vs. Wiki	9,682	15,931	547	3.320%
Google Sets vs. Amazon	9,919	809	310	27.703%

use for data generated from Google Sets as a product recommendation engine. While Amazon currently recommends products based on products purchased together (co-purchasing data), Google Sets can supplement recommendations by returning related words and phrases that users have searched for on Google. When these semantic concepts correspond to products listed on Amazon (for instance, when a word or set of words appears in the item's title), this product can then serve as a candidate recommendation for an Amazon user.

IV.3 Using Semantic Similarity for Further Refinement

Finding related words with short path lengths in the Google Sets network may return many suggestions, but a short path length alone does not guarantee that such suggestions are relevant. For example, in the "wine" to "France" graph we generated, "drink" appears immediately connected to "wine" (with a path length of 1). But despite the close proximity in the graph, users interested in products stemming from "wine" are likely not interested in every beverage available on Amazon. To prune and refine these suggestions, we can use the concept of cosine similarity, which relates semantically similar concepts by quantifying the number of shared neighbors of each of the nodes (Salton, 1975).

When we examine cosine similarity scores

for "wine" and related terms in the "wine" to "france" graphs generated from each network we studied (Table 6), we can further refine our product recommendations. From the similarity scores, it becomes clear that although "grape" may appear in the network connected to "wine", these terms are not necessarily similar. By computing cosine similarity between "wine" and every other term in the network, we can identify the most similar words, such as "beer". Therefore from our original product recommendations, we can prune many terms while maintaining such similar terms as "beer".

V. RESULTS

Here we discuss our findings and the implications they have on the interrelationships of the data sets we have examined in this paper. First, we discuss what the meaning of a shortest path indicates in the networks we studied and how networks generated from Google Sets exhibit this behavior. Next, we comment on the overlap between the three semantic networks, noting a higher overlap between Google Sets and Amazon networks than Google Sets and Wikipedia. Finally, we re-define cosine similarity and note its ability to find similarity between nodes on a deeper level than network analysis alone.

Table 6: Cosine similarity scores between "wine" and related terms in each network studied

	Google Sets	Wikipedia	Amazon
wine → france	0.0	0.0058	0.0096
wine → grape	0.0	0.02229	0.0344
wine → alcohol	0.1427	0.00921	0.0094
wine → beer	0.8007	0.1987	0.6862

V.1 Discussion

In our previous work, we showed the shortest path characteristics of the data sets that we examined throughout this paper. By running our implementation, we retrieved a path from "wine" to "france" using Google Sets. This information can be used as a supplement to a recommendation engine due to its linking of semantic concepts from a seed word to a target word.

Recommendations cannot be arbitrarily applied to any network, however. The network we choose to compare to Google Sets should share many of the same nodes, implying an overlap of the semantic structure of the network. Amazon currently has a recommendation engine that bases its recommendations on the co-purchasing data a customer has. As shown in Table 5, networks generated from Google Sets and Amazon overlap to a high degree. If Amazon were to use Google Sets as well, then the combination of the co-purchasing data along with the data available in Google Sets would allow Amazon to make more meaningful recommendations that are not based solely on co-purchasing data for that customer.

As we have seen, high network overlap and short path lengths do not always represent ideal recommendation candidates, however. To further prune the set of recommendations, we can employ semantic similarity in the form of cosine similarity, which relates nodes that share common neighbors in a network. Using cosine similarity, we can then make strong product recommendations based on data from the network generated from Google Sets.

V.2 Future Work

The techniques and algorithms and analyses we discussed in this paper can be arbitrarily expanded to any number of graphs. For example it may be interesting to study how the results in Google Sets follow the co-purchasing data of Amazon along with status updates that are made on social networks like Facebook and Twitter. These can also be linked to recent events in the news by examining graphical representations of recent articles from *The New York Times*.

It may also be possible to see how Google Sets can be applied to the discipline of psychology. The human mind can retrieve words similar to a result set like we described when given an initial word like a seed word. Further studies can be done to see how the results returned by Google Sets can be used to model the human brain and speech pathology of the human psyche.

VI. CONCLUSION

In this paper we have examined the Google Sets semantic relationship service available from Google via the Google Spreadsheet interface. We have designed and implemented a tool that will generate a graphical representation of the data available in Google Sets. We can then compare this graphical representation to graphical representations of two other data sets (Wikipedia and Amazon co-purchasing data) and run analyses to discover various characteristics common to the semantic networks. We examined which networks demonstrate small world behavior by examining their average path lengths, clustering coefficients and their power law degree distribu-

tions. The small world characteristic present in Google Sets allows us to relate concepts with short path lengths, and thus hints at the ability to recommend semantically-related concepts. Due to the high overlap in shared nodes between Google Sets and Amazon networks, we determined that Google Sets may be able to improve product recommendations typically made through Amazon co-purchasing data. Finally, we examined how semantic similarity can finalize these recommendations by narrowing the set to words and phrases that have been deemed semantically-similar, sharing many of the same neighbors in the same network.

REFERENCES

- [1] John F. Sowa and Stuart C Shapiro. *ENCYCLOPEDIA OF ARTIFICIAL INTELLIGENCE SECOND EDITION*. New Jersey: A Wiley Interscience Publication, 1992.
- [2] Mark Steyvers and Joshua B. Tenenbaum. The Large-Scale Structure of Semantic Networks: Statistical Analyses and a Model of Semantic Growth. *Cognitive Science*, 29(1):41–78, 2005.
- [3] Stanley Milgram. The Small World Problem. *Psychology Today*, 2:60–67, 1967.
- [4] Jennifer Golbeck. Generating predictive movie recommendations from trust in social networks. In *Proceedings of the 4th international conference on Trust Management, iTrust'06*, pages 93–104, Berlin, Heidelberg, 2006. Springer-Verlag.
- [5] Andrew Crossen, Jay Budzik, and Kristian J. Hammond. Flytrap: intelligent group music recommendation. In *Proceedings of the 7th international conference on Intelligent user interfaces, IUI '02*, pages 184–185, New York, NY, USA, 2002. ACM.
- [6] Reka Zsuzsanna Albert. *Statistical mechanics of complex networks*. PhD thesis, Notre Dame, IN, USA, 2001. AAI3000268.
- [7] D. J. Watts and S. H. Strogatz. Collective dynamics of 'small-world' networks. *Nature*, 393(6684):409–10, 1998.
- [8] Thomas K. Landauer, Peter W. Foltz, and Darrell Laham. An Introduction to Latent Semantic Analysis. *Discourse Processes*, (25):259–284, 1998.
- [9] G. Salton, A. Wong, and C. S. Yang. A vector space model for automatic indexing. *Commun. ACM*, 18(11):613–620, November 1975.
- [10] P. Erdős and A. Rényi. On random graphs. I. *Publ. Math. Debrecen*, 6:290–297, 1959.
- [11] Aaron Clauset, Cosma Rohilla Shalizi, and M. E. J. Newman. Power-law distributions in empirical data. *SIAM Rev.*, 51(4):661–703, November 2009.